

Group Members: Liam Puknys, Max Marcussen, Nick Nowicki, ?

Dataset: Amazon product reviews

- Source: <http://jmcauley.ucsd.edu/data/amazon/>
- Size: ~20 GB
- Total Reviews: 142.8 million

Questions:

1. What is the relationship between length of review and sentiment?
2. Star rating vs. length?
3. Star rating vs. sentiment?
4. Star rating and sentiment vs. crime rate?
5. Markov chain (put in inputs like sentiment, star rating, weather, economic state and get an example review)?
6. Are there trends in the time series of product reviews vs. trends of substitute/complement products?
7. Can large spikes in comments (+ sentiment) tell us about new product adoption?
8. Are there ever books where the majority of reviewers have not attained a reading level sufficient to read the book?
9. Books with the most profane reviews?
10. How do we determine a trustworthy review?